



Backups



Backup... in simple terms

- Having a regular fresh copy of important files on the system, which can be restored either fully or partially
- Backup can be as simple as copying files using native Linux commands like `cp`, `tar`, `rsync`, or `dump`. It can be as complex as using a third party *backup solutions*. Those can be free like Bacula, Amanda, and Mondo, or paid like NetBackup.
- The destination to which files get copied (backed up) is called *backup media*. This can be a disk, an optical disk, or a tape depending on several factors like cost and required *data retention*.
- The whole operation is governed by a *backup policy*, which entails the following:
 - The amount of data that needs to be backed up. This dictates the amount and type of needed storage.
 - The frequency by which files change. This lets you know how often backup is needed (hourly, daily, weekly...etc.)
 - The duration for which data needs to be kept (data retention). This, too, controls the amount and type of backup media.

LAB: Calculate the amount of data that needs to be backed up


- Depending on your environment, you'll probably have a number of machines on your network that needs backup. In this lab you write a script that will calculate the total amount of space and the used one on a number of systems:

```
#!/bin/bash
TOTAL=0
USED=0
OUT=""
MACHINES=( 192.168.0.103 192.168.0.101 )
for m in "${MACHINES[@]}"
do
    OUT="$OUT~`ssh $m 'df / | tail -n +2'`"
done
LENGTH=${#MACHINES[@]}
LENGTH=$((LENGTH + 1))
for (( i=2; i<=$LENGTH; i++ ))
do
    USED_TEMP=`echo $OUT | cut -d "~" -f $i | cut -d " " -f 3`
    TOTAL_TEMP=`echo $OUT | cut -d "~" -f $i | cut -d " " -f 2`
    USED=$((USED+USED_TEMP))
    TOTAL=$((TOTAL+TOTAL_TEMP))
done
TOTAL=$((TOTAL/1024/1024))
USED=$((USED / 1024 / 1024))
echo "Total is " $TOTAL " GB"
echo "Total used is " $USED " GB"
```

LAB: Calculate the frequency by which files change

- In this lab, we are going to write a script that sums up the sizes of files that were changed in the last 24 hours. The results are mailed to root. Such a script can be added to a cron job that will carry it out periodically.

```
#!/bin/bash
SUM=0
FILES=`find / -mount -type f -mtime -1 | xargs du -k | awk
'{print $1}'`
for i in $FILES; do
    SUM=$((SUM+i))
done
SUM=$((SUM/1024))
echo "$SUM MB" | mail -s "$HOSTNAME free space" root
```



Using nc for non-SSL-based remote backups

- ▶ Short for netcat, a utility that is used to make a lot of network-related operations. We are using it here as a daemon to aid in making remote backups.
- ▶ Data is transferred unencrypted so nc should not be used to backup sensitive data
- ▶ On the destination server, we used netcat to receive files by choosing a port and letting nc wait for input
`nc -l 20000 | tar -xvf -`
- ▶ Tar can be replaced with any other command that can read from standard input (like dump)
- ▶ On the client machine, we can issue a command like this
`tar -cvf - file(s) | nc destination_ip 20000`



Using cp

- It is the simplest native tool used for backup
- It can be used to copy files and directories to a local or a remote directory.
- Typical usage is as follows:
`cp -ax source destination`
where *a* is *archive mode*, which preserves permissions and other file attributes. The *x* switch ensures that the command does not leave the filesystem to other mounted filesystems
- To include directories
`cp -Rax source destination`
- Shortcomings of using `cp`:
 - Does not support resuming broken operations
 - Does not use SSL (when copying to remote directories)
 - The remote destination directory must be shared through SAMBA or NFS, which provides an extra burden
 - Does not support differential backups
- Used for very basic backup operations

Using tar

- Originally designed for tape backups (short for Tape Archive)
- Can still be used to backup to tapes

```
tar -cf /dev/rmt0 /root
```
- Can also be used to backup to a regular file (called tarball)

```
tar -cf /backup/mybackup.tar /root
```
- Commonly used with gzip to compress tarballs

```
gzip /backup/mybackup.tar
```
- Use the x argument to restore a number of files or the whole archive

```
tar -xf /backup/mybackup.tar [file(s)]
```
- It can be used with remote directories and it supports SSL (using SSH)

```
tar -cvf - file(s) | (ssh remote_server 'tar -C directory -xvf -')
```
- The tar method has the following shortcomings:
 - Does not support resuming broken operations
 - Does not support incremental/differential backup



Using rsync

- A robust command used to copy files locally or to a remote server
- Supports SSL natively
- Can resume broken operations
- Supports differential/incremental backups because only new or changed files get copied
- Has options to synchronize directories. This is very useful when uploading files that change a lot in mirrored locations like webserver production and development machines
- When used in backup, it is most commonly used with the -a command switch (for archive), which preserves file permissions, symbolic links, and other file attributes
- Can be used with -z to compress files while being copied. This slightly speeds up backups on networks but consumes CPU
- Can be used with -P to provide a progress bar.

Using dump

- Sometimes it is considered the standard backup tool for Linux/UNIX
- It provides ten levels of backup (0-9), where 0 is full backup, every integer above 0 makes dump examine any new or changed files since that last integer backup. For example level 5 will take all changed/added files since level 4.
- It can be used to backup to tapes. When the tape is full, it will continue the operation when the next tape is ready.
- It can also be used to backup to a file, and – combined with SSH – can be used to backup to remote directories
- The typical usage for backup
`dump 0f /dev/rmt0 /filesystem`
where 0 is the level, /dev/rmt0 is the backup media (can be a regular file), and /filesystem is the path to the filesystem or directory to be backed up
- When restoring (must be inside the directory where you want to restore)
`restore rf /path/to/mybackup.dmp`



LAB: cloning an entire Linux machine using dump and dd

- This method is used to backup an entire machine to a disk so that it can be used to boot to an identical copy, or to restore the original machine in case of disaster. The following procedure does not assume that the two machines are having identical disk sizes
- The operation can be summarized as follows:
 - [Source Machine]: Turn off and boot from the installation DVD, enter "rescue mode". You can skip this step if are absolutely sure no files will be changed during the cloning process
 - [Destination machine]
 - Start the destination machine using the installation DVD and enter "rescue mode"
 - Choose to start network manager
 - Choose not to mount the disks

- Clone the superblock from the source machine to the destination one using dd and nc:

- [Destination machine] : `nc -l 20000 | dd of=/dev/sda`

- [Source machine]: `dd if=/dev/sda bs=446 count=1 | nc 192.168.0.105 20000`

- We used 446 because it contains the bootstrap information. The remaining bytes are 64 for partition table (we created a different partition table), and 2 for signature totaling 512 (boot block).

- [Destination machine]:

- Using fdisk, create a primary partition for boot

- Create a new partition for the volume group (assuming the root volume group is using only one physical volume)

- Create a volume group with the same name as the one on the source machine

- Create all the required logical volumes that were on the source machine. Sizes do not have to be the same

- Don't forget to create an lv for the SWAP space.

- Using mkfs.ext4, create new filesystems on the new logical volumes. Use mkswap to format the swap space.

- Mount the logical volumes on temporary directories

- Start cloning files from source to destination machines:
 - [Destination machine while in /mnt/slash]
`nc -l 20000 | restore -rf -`
 - [Source machine] `dump 0f - /dev/vg_redhat01/lv_root | nc 192.168.0.5 20000`
- [Destination machine] After all files are cloned, examine the fstab file and correct any problems that might prevent the system from booting. For example having the original UUID of /dev/sda1, which has to be changed or replaced directly with /dev/sda
- [Destination machine] reboot the machine using the installation DVD and enter rescue mode. This time you can opt to let it mount the disks for you under /mnt/sysimage.
- [Destination machine] make /mnt/sysimage your root filesystem using `chroot /mnt/sysimage`
- [Destination machine] install GRUB using the following commands:
 - `$ grub`
`grub> root (hd0,0)`
`grub> setup (hd0)`
`grub> quit`
- Reboot the system and ensure that it is running as expected